

Parallel Python

Interactive parallel computing with IPython

Monte Lunaceck and Thomas Hauser @colorado.edu

Why Python?

- Relatively easy to learn
- Interactive
- Cross platform
 - Linux
 - Windows
 - OS/X
- Large number of packages
 - Numpy – for numerical computation
 - Matplotlib - plotting
 - Ipython – interactive shell

IPython

- Supports different styles of parallelism
 - *Single program, multiple data (SPMD) parallelism*
 - Multiple program, multiple data (MPMD) parallelism
 - Message passing using MPI
 - *Task farming*
 - Combinations of above approaches
 - Custom user defined approaches
- Interactive approach to
 - Development
 - Execution
 - Debugging
 - Monitoring

Overview

- How it works
- Direct interface
 - Executing remote commands
 - Data management
 - Blocking, non-blocking
- Task interface
 - Compare with pbsdsh
 - Simple, robust, and load-balanced


How it works

Architecture

- Engine
 - Takes Python commands over a network connection
 - One engine per core
 - Blocks code: controller has an asynchronous API
- Controller client
 - Interface for working with a set of engines
 - Hub
 - Collection of schedulers
 - View objects represent a subset of engines
 - A **Direct** interface, where engines are addressed explicitly.
 - A **LoadBalanced** interface, where the Scheduler is trusted with assigning work to appropriate engines.

Parallel Architecture

- Hub
 - keeps track of engine connections, schedulers, clients
 - Cluster state
- Scheduler
 - All actions that can be performed on the engine go through a scheduler



Start your Engines

- Simplest method
 - `ipcluster start --n=4`
 - MPIEXEC/MPIRUN mode
 - SSH mode
 - PBS mode for batch processing
- Other methods
 - Ipcontroller
 - ipengine
- Don't forget to turn off your engine!

Use in HPC environments

- `ipython profile create --parallel -profile=mpi`

Generating default config file: u'/home/molu8455/.ipython/profile_mpi/ipython_config.py'Generating default config file: u'/home/molu8455/.ipython/profile_mpi/ipython_qtconsole_config.py'Generating default config file: u'/home/molu8455/.ipython/profile_mpi/ipcontroller_config.py'Generating default config file: u'/home/molu8455/.ipython/profile_mpi/ipengine_config.py'Generating default config file: u'/home/molu8455/.ipython/profile_mpi/ipcluster_config.py'Generating default config file: u'/home/molu8455/.ipython/profile_mpi/iplogger_config.py'

- Edit the file `ipcluster_config.py`
 - `c.IPClusterEngines.engine_launcher_class = 'MPIEngineSetLauncher'`

Steps

- Request a resource with qsub:
 - Interactive
 - Batch
- Start the engines
 - `ipcluster start --profile=mpi &`
- Parallel python work ...
- Stop the engines
 - `ipcluster stop --profile=mpi`

Example in a PBS setting

```
#PBS -N example
```

```
#PBS -j oe
```

```
#PBS -l walltime=00:45:00
```

```
#PBS -l nodes=2:ppn=12
```

```
#PBS -q janus-debug
```

```
# any config relevant to your local environment here
```

```
# switch to the PBS working directory
```

```
cd $PBS_O_WORKDIR
```

```
# Start the IPython cluster and wait a little for engines to come up
```

```
ipcluster start --profile=mpi &
```

```
sleep 90
```

```
# Run your own code that uses the ipcluster
```

```
python work.py
```

```
# Stop the cluster when done
```

```
ipcluster stop --profile=mpi
```

A few simple concepts

- The **client**: lightweight handle on all engines of a cluster

```
from IPython.parallel import Client  
rc = Client(profile='mpi')
```

- The **views**: “slice” the client with specific execution semantics

- DirectView: direct execution on *all* engines

```
dview = rc.direct_view()
```

- LoadBalancedView: run on *any one* engine

```
lview = rc.load_balanced_view()
```

Direct Interface

Direct Interface

- capabilities of each engine are directly and explicitly exposed to the user
- Blocking and non-blocking
 - `AsyncResult`
- Remote execution
 - `apply`
 - `execute`
 - `map`
 - Remote function decorators
- Managing data
 - `push`, `pull`
 - `scatter`, `gather`

Blocking vs. Non-blocking

- Blocking
 - waits until all engines are done executing the command
`dview.block=True`
- Non-blocking
 - returns an `AsyncResult` immediately
 - Get data later with `get` method
 - Check the status with `ready`
 - Wait
`dview.block=False`

Apply Example

- Serial

```
import os
pid = os.getpid()
print pid
```
- Parallel

```
with dview.sync_imports():
import os
dview.block=True
pids = dview.apply(os.getpid)
dview.block=False
ar = dview.apply(os.getpid)
dview.wait(ar)
pids = ar.get()
```


Apply

- For convenience
 - `apply_sync`
(`block=True`)
 - `apply_async`
(`block=False`)

```
def wait(t):  
    import time  
    tic = time.time()  
    time.sleep(t)  
    return time.time()-tic  
  
ar = dview.apply_async(wait,1)  
if ar.ready() == True:  
    print "ready"  
else:  
    print "not-ready"  
dview.wait(ar)
```

Execute

- Commands can be executed as strings on specific engines

```
dview['a'] = 10
dview['b'] = 20
dview.execute('c=a*b')
rc[1].execute('c=a+b')
print dview['c']
```

[200, 30, 200, 200, 200, 200, 200, 200, 200, 200, 200, 200]

Map

- Applies a sequence of to a function element-by-element
- Does not do dynamic load balancing
- `map_sync`, `map_async`

```
def fsquare(x):  
    return x**2  
xvals = range(32)  
s_result = map(fsquare, xvals)  
dview.block=True  
p_result = dview.map(fsquare, xvals)
```

Push and Pull

- Limited to dictionary types

```
dview.block=True
values = dict(a=1.03234,b=3453)
dview.push(values)
a_values = dview.pull('a')
print a_values
dview.push(dict(c='speed'))
print dview.pull('c')
dview['c'] = 'speed'
print dview.pull('c')
```

Scatter and Gather

- Partition a sequence and push/pull

```
dview.scatter('a', range(24))
```

```
print dview['a']
```

```
all_a = dview.gather('a')
```

```
print all_a
```

```
# Use
```

```
dview.scatter('x', range(64))
```

```
dview.execute('y=map(fsquare, x)')
```

```
y = dview.gather('y')
```

```
print y
```

Task Interface

Task interface

- Presents the engines as a *fault tolerant, dynamic load-balanced* system of workers
- No direct access to individual engines
- Simple and powerful

```
lview = rc.load_balanced_view()
```

- Components:
 - map
 - Parallel function decoders
 - Dependencies
 - functional
 - graph

Example: parameter study

```
#!/bin/bash
#PBS -N example_mn_1
#PBS -q janus-debug
#PBS -l walltime=0:01:30
#PBS -l nodes=2:ppn=12
```

```
cd $PBS_O_WORKDIR
```

```
pbsdsh
$PBS_O_WORKDIR/wrapper.sh 0
```



```
#!/bin/bash
PATH=$PBS_O_WORKDIR:$PBS_O_PATH
TRIAL=$(( $PBS_VNODENUM + $1 ))
simulator $TRIAL > $PBS_O_WORKDIR/sim.$TRIAL
```



```
#!/bin/bash
#PBS -N example_mn_3
#PBS -q janus-debug
#PBS -l walltime=0:01:30
#PBS -l nodes=2:ppn=12

cd $PBS_O_WORKDIR


count=0
np=`wc -l < $PBS_NODEFILE`

for i in {1..5}
do
    pbsdsh
    $PBS_O_WORKDIR/wrapper.sh $count
    count=$(( $count + $np))
done
```


Problems

- Core failure to run
- Load-balance

Weaknesses of batch




time →




time ↘

Load balancing



time →



time ↘

simulator.py

```
def imulator(x):
    import subprocess
    cmd = "./simulator " + str(x) + " 2 4 > trial." + str(x)
    p = subprocess.Popen(cmd, shell=True, stdout=subprocess.PIPE)
    (err,out)=p.communicate()
    status = p.returncode
    return "%s, %s" % (out,err)

print simulator(1)

from IPython.parallel import Client
rc = Client(profile='mpi')
lview = rc.load_balanced_view()
r = lview.map_async(simulator, range(100))
print r.get()
```

```
#!/bin/bash
#PBS -N example_mn_3
#PBS -q janus-debug
#PBS -l walltime=0:01:30
#PBS -l nodes=2:ppn=12
```

```
cd $PBS_O_WORKDIR
```

```
ipcluster start --profile=mpi &
sleep 90
```

```
python simulator.py
```

```
ipcluster stop --profile=mpi
```

Interactive MPI

- Create MPI functions

```
def psum(a):  
    s = np.sum(a)  
    rcvBuf = np.array(0.0, 'd')  
    MPI.COMM_WORLD.Allreduce([s, MPI.DOUBLE], [rcvBuf, MPI.DOUBLE], op=MPI.SUM)  
    return rcvBuf
```

- Call interactively from direct view

Conclusions

- Direct interface
 - Quickly parallelize python code
 - Less control than MPI
- Task interface
 - Fault-tolerant, load-balanced, simple
- MPI interactive
 - Great for developing, debugging
- Tutorials:
 - <http://ipython.org/ipython-doc/dev>
 - <http://minrk.github.com/scipy-tutorial-2011>