

Python and Parallel computing: an overview

Fernando Pérez

`Fernando.Perez@berkeley.edu`

Helen Wills Neuroscience Institute, UC Berkeley

Python Computing for Science - AY250

Outline

- 1 HPC and parallelism
- 2 Python
- 3 Numpy and Scipy

Outline

- 1 HPC and parallelism
- 2 Python
- 3 Numpy and Scipy

Data explosion

Our data sets are getting **huge!**

- Sloan Digital Sky Survey Data Release 7 : ~65 Terabytes
- Neuroimaging: 1 hour \rightarrow a few gigabytes
- etc...

A simple problem: solve $Ax = b$ via Gaussian elimination

- For A an $n \times n$ matrix, cost is $\approx \frac{1}{3}n^3$ floating point operations
- On a computer with ~100MFlops sustained performance

n	time
10	$\mathcal{O}(1) \mu s$
100	$\mathcal{O}(1) ms$
1000	$\mathcal{O}(1) s$
10000	$\mathcal{O}(1) hour$
100000	$\mathcal{O}(1) month$
1000000	$\mathcal{O}(1) century$

Why do we need fast algorithms?

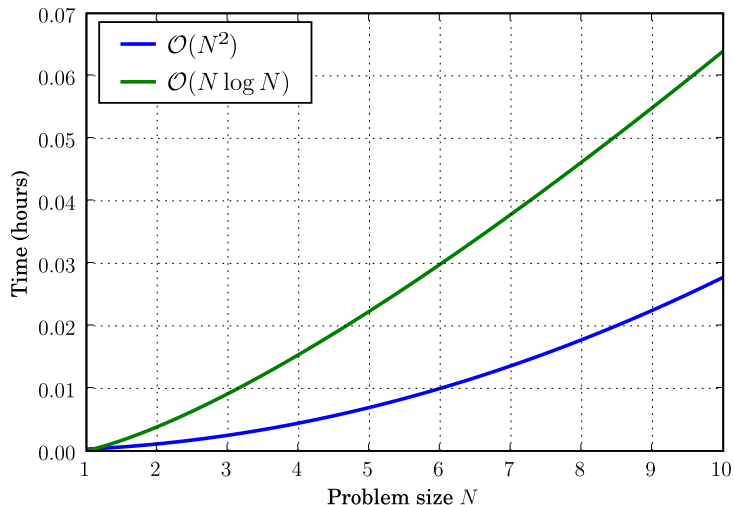
Because computers are getting bigger and faster!

Why do we need fast algorithms?

Because computers are getting **bigger and faster!**


Why do we need fast algorithms?

Because computers are getting **bigger and faster!**




Why do we need fast algorithms?

Because computers are getting **bigger and faster!**




Why do we need fast algorithms?

Because computers are getting **bigger and faster!**



Parallel computing: why should we care?


Because reality looks like this:



Sources: Intel, Microsoft (Sutter), Stanford (Olukotun, Hammond) & Berkeley (Yelick)

Parallel computing: why should we care?


Because reality looks like this:



Sources: Intel, Microsoft (Sutter), Stanford (Olukotun, Hammond) & Berkeley (Yelick)

We can't escape thermodynamics

Moore's Law Extrapolation: Power Density for Leading Edge Microprocessors



Power Density Becomes Too High to Cool Chips Inexpensively


Sources: Shekhar Borkar, Intel Corp & Kathy Yelick, UC Berkeley

The vendor's solutions:

- Multicore chips: even in your laptop.
- Graphics cards for general computing: > 128 'processors' per card.
- Clusters, cloud, ...

We can't escape thermodynamics

Moore's Law Extrapolation:
Power Density for Leading Edge Microprocessors



Power Density Becomes Too High to Cool Chips Inexpensively

Sources: Shekhar Borkar, Intel Corp & Kathy Yelick, UC Berkeley

The vendor's solutions:

- Multicore chips: even in your laptop.
- Graphics cards for general computing: > 128 'processors' per card.
- Clusters, cloud, ...

How fast can we go?

Exercise: Derive and plot Amdahl's law

Upper bound on speedup achievable via parallelization

- s : serial fraction of total work to be done
- $1 - s$: parallelizable fraction
- p : number of processors used
- Let T_n be the time for a job with n processors in use.
- Derive the possible speedup T_1 / T_p .
- Compare to a simple function of s .

Amdahl's law: a logical limit

Consider solving a problem where communication has zero cost

- Using 1 processor: $T_1 = s + (1 - s) = 1$
- Using p processors: $T_p = s + \frac{1-s}{p}$

The total speedup possible:

$$\frac{T_1}{T_p} = \frac{1}{s + \frac{1-s}{p}} < \frac{1}{s}$$

Lesson: The serial fraction is a hard limit!

Amdahl's law: a logical limit

Consider solving a problem where communication has zero cost

- Using 1 processor: $T_1 = s + (1 - s) = 1$
- Using p processors: $T_p = s + \frac{1-s}{p}$

The total speedup possible:


$$\frac{T_1}{T_p} = \frac{1}{s + \frac{1-s}{p}} < \frac{1}{s}$$

Lesson: The serial fraction is a hard limit!

Plot Amdahl's law

```
def amdahl(s, p):  
    return 1.0/(s+(1.0-s)/p)  
  
p = np.logspace(0, 16, 9, base=2)  
ax = plt.subplot(111)  
  
for s in [0.5, 0.25, 0.1, 0.05]:  
    sp = amdahl(s, p)  
    ax.semilogx(p, sp, '-o', label='s=%.2f' % s, basex=2.0)  
  
ax.set_xlabel('Processors')  
ax.set_ylabel('Speedup')  
ax.set_title('Speedup as a function of serial fraction')  
ax.legend()  
ax.grid()  
  
plt.show()
```

Amdahl's law



General Principles of Parallel Programmig

Credit: Kathy Yelick, UCB.

- Finding enough parallelism (Amdahl's law)
- Granularity: bite-sized chunks for each unit...
 - But need large enough amount of work to hide the overhead
- Locality
 - large memories are slow, fast memories are small.
- Load balance
- Coordination and synchronization
 - Communication is expensive...
 - But getting the wrong answer fast doesn't cut it.
- Model performance, profile, profile, profile...
 - If you didn't measure it, you don't actually know.
 - Your intuition is wrong.

The thirteen dwarves: patterns in parallel programming

The Landscape of Parallel Computing Research: A View from Berkeley

- 1 Dense Linear Algebra
- 2 Sparse Linear Algebra
- 3 Spectral Methods
- 4 N-Body Methods
- 5 Structured Grids
- 6 Unstructured Grids
- 7 MapReduce
- 8 Combinational Logic
- 9 Graph Traversal
- 10 Dynamic Programming
- 11 Backtrack and Branch-and-Bound
- 12 Graphical Models
- 13 Finite State Machines

http://view.eecs.berkeley.edu/wiki/Dwarf_Mine

Outline

- 1 HPC and parallelism
- 2 Python
- 3 Numpy and Scipy

Threading and parallelism in Python: overview

- Multiple implementations of the Virtual Machine:
 - CPython: pure C, 'reference'
 - IronPython: .NET
 - Jython: Java
- Their threading behaviors differ, I'll focus on CPython
- Native threads supported, but of **limited use**.
- **Global interpreter lock** (GIL): only **one** thread can modify any python data structure
- **No language-specific primitives** for parallelism.

Parallelism in Python

• In-process (mind the GIL)

- [Data parallelism](#) with threaded libraries
- Numpy/scipy can use a [threaded ATLAS](#)
- [Numexpr](#): a 'numpy VM'
- [Theano](#): a library that thinks it's a compiler
- GPU-based solutions: [PyCuda](#)/[PyOpenCL](#), [scikits.cuda](#).
- Hand-written threaded code...

• Out-of-process

- The [multiprocessing](#) module
- Python [futures](#) (only in Python 3.2 or newer)
- Communicating Sequential Processes, ParallelPython, ... many more
- [IPython](#) (I'm obviously biased). For more on IPython, see:
<https://github.com/ipython/ipython-in-depth>.

Multiprocessing

Module: multiprocessing

- Built-in since version 2.6 (available for earlier versions)
- An API that closely follows the [threading](#) API, but using processes
- Useful high-level objects
 - Process, Process pool, Namespaces, Listeners, ...
- Uses `fork()` on posix (hence there are some **limitations**)

A simple example

```
from multiprocessing import Process

def f(name):
    print 'hello', name

if __name__ == '__main__':
    p = Process(target=f, args=('bob',))
    p.start()
    p.join()
```


Outline

- 1 HPC and parallelism
- 2 Python
- 3 Numpy and Scipy

Numpy and Scipy: 'Out of the box' parallelism?

- **Not great...**
- Can be built against a threaded ATLAS or the Intel Math Kernel Library (MKL)
 - This can give multithreaded support to many linear algebra operations.
- Manual effort with C/Fortran + OpenMP can give you some gains...
 - but with a fair amount of pain

Numexpr

An expression compiler for numpy

Approach

- Compile Numpy expressions to equivalent Python code...
- Block operations carefully
- execute on a special-purpose mini-VM (written in C)

Benefits

- Reduce the use of temporaries.
- Be cache-friendly.
- Support threads natively for all operations.
- Support Intel Vector Math Library and MKL.

Numexpr usage

Evaluating simple expressions

```
>>> import numpy as np
>>> import numexpr as ne

>>> a = np.arange(1e6)    # Choose large arrays for high
    performance
>>> b = np.arange(1e6)

>>> ne.evaluate("a + 1")  # a simple expression
array([ 1.00000000e+00,  2.00000000e+00,  3.00000000e+00, ...,
        9.99998000e+05,  9.99999000e+05,  1.00000000e+06])

>>> ne.evaluate('a*b-4.1*a > 2.5*b')  # a more complex one
array([False, False, False, ..., True, True, True], dtype=bool)
```

Comparisons to Numpy and thread usage

```
>>> timeit a**2 + b**2 + 2*a*b
10 loops, best of 3: 35.9 ms per loop
```

```
>>> ne.set_num_threads(1) # use 1 thread (on a 6-core machine)
```

```
>>> timeit ne.evaluate("a**2 + b**2 + 2*a*b")
100 loops, best of 3: 9.28 ms per loop # 3.9x faster than NumPy
```

```
>>> ne.set_num_threads(4) # use 4 threads (on a 6-core machine)
```

```
>>> timeit ne.evaluate("a**2 + b**2 + 2*a*b")
100 loops, best of 3: 4.17 ms per loop # 8.6x faster than NumPy
```